Parsing Foundations

A Study of Data and Parsing

Unfinished Draft of First Edition

Parsing Foundations

A Study of Data and Parsing

Unfinished Draft of First Edition

Richard R. Masters
Seattle, WA

This book was typeset using LaTeX software.

Preface

Parsing is an important and well-established area within the field of software development and nearly every programmer can benefit from a solid understanding of parsing.

There is an enormous amount of information on the topic of parsing. However, it requires a serious investment of time (and often times some money) to identify, locate, acquire, organize, and understand the books, papers, and online resources that will help you gain a solid understanding of the relevant issues. This book was written to help the reader navigate this large body of information. There is also a significant amount of useful research which has been around for several decades but has been hidden in obscurity, and so a goal of this book is to present this research to a wider audience.

This book does not assume that the reader is an expert with set-builder notation[1] or category theory or that the reader is familiar with specific mathematical symbols that have been used in previous papers on the subject. Thorough explanations are preferred over terse explanations.

This book starts with a thorough explanation of data because a good understanding of parsing rests on a good understanding of data, which is what a parser acts upon.

This document is licensed under the Creative Commons Zero v1.0 Universal license[2]. That means you are free to copy this content, distribute it, include it in your own content, or sell it, and you may do so *without attribution*. Just copy what you want and move on.

Audience

This book can serve as an introduction to topics on data and parsing for the working software developer or student of computer science. The book covers foundational topics of parsing using concrete examples and common tools used in the field. In an academic setting, the book would serve well as a companion book for courses that only cover parsing as part of a larger topic, such as compiling.

Prior Work

Prior to the writing of this book there were no public domain books on parsing.

While there are quite a few papers and articles written for academic journals, these articles are typically written for a narrow audience by use of vocabulary and notation. Also, many important papers on parsing are not freely available; one must purchase a copy in order to read it, or have access to a very good library. Finally, while there are some good online resources such as the work of Federico Tomassetti et al, they are lost within a sea of information that is either outdated or poorly produced, or only covers a limited set of topics.

The following is a list of commercially published books that the author has been able to review.

1. Dick Grune, Ceriel J.H. Jacobs
   Parsing Techniques, A Practical Guide, Second Edition, 2008

   This is the best book focused on parsing for the general audience.

2. Alfred V. Aho, Monica S. Lam, Ravi Sethi, Jeffrey D. Ullman
   Compilers: Principles, Techniques, and Tools, Second Edition, 2007

   This is a very popular book that covers lexing, parsing, and compilation. The chapters on lexing are particularly good, covering DFA construction. This book is commonly known as the "Dragon Book" due to the illustration on the cover.

---

[1] https://en.wikipedia.org/wiki/Set-builder_notation
[2] https://creativecommons.org/publicdomain/zero/1.0/legalcode

3. S. Sippu, E. Soisalon-Soininen
   Parsing Theory, 1988

   This book describes parsing in terms of category theory and is appropriate for students who have taken graduate level classes in mathematics.

4. Nigel P. Chapman
   LR Parsing, Theory and Practice, 1987

   This book describes parsing in the language of mathematics and is appropriate for students who have taken undergraduate level classes in mathematics.

5. John-Paul Tremblay, Paul G. Sorenson
   The Theory and Practice of Compiler Writing, 1985

6. William A. Barrett, John D. Couch
   Compiler Construction: Theory and Practice, 1979

Terms and Symbols

Many of the ideas covered in this book originated in the 1950s and 1960s, near the dawn of the computer age. In the decades since, researchers and educators have defined quite a few names for various concepts surrounding data and parsing. They have also adopted numerous symbols and notations as a convenient shorthand to write down their ideas. While it can be tedious to learn (and explain) these names and symbols, it is necessary if you want to be able to follow discussions or read original research papers on these topics. In some cases, words that you may be familiar with in common language may have a different definition in academic papers. This can be confusing if you are not already familiar with that term. Therefore, in this book, we will error on the side of explaining the academic terms and symbols when they are used. Sometimes it takes a number of definitions conspiring together to explain a larger concept, so bear with us if it seems like we are just defining words for no apparent reason.

Prerequisites

For the chapters on data, the reader should have a basic knowledge of computers, such as knowing about a monitor, mouse, disk, file, and an application.

Table of Contents

Chapter 1

Parsing Concepts

1.1 Strings

1.2 Languages

1.3 Grammars

1.3.1 Terminals and Tokens

Terminals are typically referred to as tokens in practice. Technically, a terminal is just a language element, while a token typically contains additional information about the context in which the element was recognized.

1.3.2 Nonterminals

Nonterminals are often referred to as variables.

1.3.3 Productions

A production is a rule that map a nonterminal to a sequence of terminals.

1.4 Parse Trees

1.4.1 Concrete Syntax Tree

A tree that accurately reflects the grammar that produced it.

1.4.2 Abstract Syntax Tree

This is a parse tree that is abstraction of a concrete (or grammar-based) parse tree in order to be independent of the grammar that produced it.

1.4.3 Document Parsing

```
parse_tree = parse(input)
```

1.4.4 Stream Parsing

```
for event in parse(input):
    handle_event(event)
```

For bottom up parsing, events are produced as wholly formed subtrees are recognized. For top down parsing, events are produced for creation of the root node, and for children of a previously parsed node.

1.4.5   Incremental Parsing

Parse trees are needed for incremental parsing:

```
reparse(old_parse_tree, text_delta)
```

`https://channel9.msdn.com/Blogs/Seth-Juarez/Anders-Hejlsberg-on-Modern-Compiler-Construction`

1.5   Parsing Applications

There are important differences between parsing data files versus parsing computer languages.

1.6   Parsing Software

There are handwritten and generated parsers.

Chapter 2

Lexers

2.1 The Role of the Lexer

Lexemes and Tokens

Why do lexing independently from parsing?

Wirth, in Compiler Construction, states

"A partitioning of the compilation process into as many parts as possible was the predominant technique until about 1980, because until then the available store was too small to accommodate the entire compiler"

A pipeline supports separation of concerns and allows reuse of a lexer with other parsers that accept a stream of tokens.

```
https://rosettacode.org/wiki/Compiler/lexical_analyzer
https://rosettacode.org/wiki/Compiler/Sample_programs
```

2.1.1 Lexers, Tokenizers, and Scanners

The simplistic view is these are all the same. The more nuanced view is that there are distinctions between these words that can be helpful.

A *lexer* is software that takes data as input and outputs tokens, (i.e. it acts as a tokenizer) using a scanner to process the input data.

2.1.2 Scannerless Parsing

The concept of *scannerless parsing* was formalized by Salomon and Cormack in "Scannerless nslr(1) parsing of programming languages" July 1989 It may not be immediately clear why simply *not* using a scanner is an idea being discussed in the year 1989. If you read Knuth's 1965 paper on LR parsing, there is no mention of scanning. One would presume that the lack of a scanner was a normal thing. When Knuth discusses parsing a string such as *abcde* it his paper, he presumes each of those letters is a separate terminal. Apparently, there is no scanner. But strictly speaking, his paper was simply agnostic about the nature of the terminals such as how they were created or whether they "contain" information or structure, such as a sequence of characters. The parsing process he describes presumes the input string is composed of a pure sequence of atomic elements.

In practice, however, a scanner was typically used because it allows the resulting grammar to be much simpler than a grammar in which every character is a terminal. A simpler grammar can be parsed with a simpler parser, which is an essential benefit of using a scanner.

2.2   Scanning

2.2.1   Input Data

2.2.2   Rules for Tokens

2.2.3   Literals

2.2.4   Regular Expressions

2.2.5   Pseudo-Tokens

2.2.6   Context-Free and Context-Aware Scanning

2.3   Scanner Conflicts

When two different token rules match the same string, that is a *scanner conflict*.

2.3.1   Traditional Lexer Precedence

A *traditional lexer* resolves conflicts between tokens using *traditional lexical precedence* as defined by [Denny, Def 2.2.6] and repeated here.

Traditional lexers have been context-free lexers. They also resolve conflicts using the *longest match* rule first, followed by *positional priority* as described in the next two sections, respectively.

2.3.1.1   Longest Match Rule

2.3.1.2   Positional Priority

2.3.2   Longest Match Conflict

A longest match conflict occurs when a token chosen by a traditional lexer leads to an error state in the lexer or parser, whereas if a different token, such as a shorter one was chosen, the parse would succeed.

For example, 1..10

If 1 and 1. are valid tokens (for integers and floating point numbers), then these sequences are valid:

```
(1.) (.10)
(1) (..) (10)
```

Longest match favors the first sequence, but that token sequence is invalid to the parser, but the second sequence is valid.

The presence of this problem is a combination of:

- The traditional lexical and grammatical model cannot express this language without special handling.

- Trying to use or design a language that requires this special handling

For example, scannerless parsing with an LR parsing algorithm will be able to parse this string. Therefore, one might argue that the burden is on the parsing software to handle this issue. The counterargument is that the language designer can avoid this issue by being thoughtful of it during the design process and verifying that these conflicts do not exist using proper tools.

(Check to see whether lex's lookahead operator can solve this. this would work if the length used for longest match rule includes the lookahead string (seems unlikely but who knows))

Assuming you are using traditional scanning software, a strategy for handling longest match conflicts, which makes the grammar less portable between various parsing tools.

2.3.3   Context-Invasive Token Conflicts

A scanner conflict that only occurs because a token is active that is not associated with the current parsing context is given the name *context-invasive conflict* in this book.

2.3.4 Resolving Lexical Conflicts

2.3.4.1 Lexical Modes

Lexical modes are a tool for resolving lexical conflicts by creating separate sets of active tokens.

2.3.4.2 Lexical Lookahead Operator

A lexical lookahead operator can specify a required pattern to follow a lexeme in order to match the token. This pattern is not included in the lexeme.

2.3.4.3 Lexical Predicates

A lexical predicate is arbitrary code that runs in order to determine whether to match the token.

2.3.4.4 Match Multiple Tokens

2.3.5 Academic History of Token Conflicts

Token conflicts were first formalized by Nawrocki in 1991.
  Keynes:

  universal quantified predicate (text p. 5, pdf p. 11)
`http://www.open-std.org/jtc1/sc22/open/n3187.pdf`

  set comprehension term p. 20
`https://www.cs.ox.ac.uk/files/3389/PRG68.pdf`

  Denny

Chapter 3

Parsing Algorithms

3.1   Left-to-Right Parsing

3.2   Left and Right Derivations

3.3   Top-down Parsing

This produces a left-most derivation.

3.4   Bottom-up Parsing

This produces a right-most derivation.

3.5   Parsing Families

3.6   The LL Parsing Family

3.7   The LR Parsing Family

3.8   CYK Parsing

3.9   Early Parsing

3.10   PEG Parsing

3.11   Parser Combinators

Chapter 4

LL Parsing

4.1 LL Table Structure

4.2 LL Table Building

4.3 ALL Parsing

4.4 GLL Parsing

Chapter 5

LR Parsing

5.1 LR Table Structure

5.2 LR Table Building

5.2.1 LR Table Conflicts

5.2.2 Avoiding LR Table Conflicts

5.2.3 Resolving LR Table Conflicts

5.3 SLR(0) Parsing

5.4 Canonical LR(1) Parsing

Uncompressed, not large. "Large" is a carryover from a time when an LR table challenged memory resources of a typical computer. The difference now, between 100K and 1M is not much. In the 1980's, that could have meant much more.

5.5 LALR Parsing

5.6 Minimal LR Parsing

5.7 GLR Parsing

Chapter 6

Indentation-Sensitive Parsing

6.1 Indentation-Sensitive Languages

Python, YAML, and Haskell are indentation sensitive.

6.2 Parsing Indented Blocks

6.2.1 Grammar Annotations

6.2.1.1 Grammar terms annotated with Counts

Indentation Sensitive Languages
2006 Leonhard Brunauer, Bernhard Muhlbacher
`https://www.google.com/search?q=LeonhardBrunauer+and+Bernhard+M%C3%BChlbacher.+Indentation+sensitive+languages`

The scheme allows a specific character to be matched a specified number of times.

For example, a tab character, represented by $\rightarrow$, can be repeated m times in a grammar with this syntax:
$C_{\rightarrow}^{m}$

This rule states the start variable expands to zero tabs followed by a statement:
$S \rightarrow C_{\rightarrow}^{0} Statement$

This rule states that the statement after an "if" condition has greater indentation:
$C_{\rightarrow}^{m} Statement \rightarrow C_{\rightarrow}^{m} "if" \, Condition \, ":" \, newline \, C_{\rightarrow}^{m+1} Statement$

```
if true:
    a=1
```

6.2.1.2 Relationship Annotated BNF

Principled Parsing for Indentation-Sensitive Languages 2013
`https://michaeldadams.org/papers/layout_parsing/LayoutParsing.pdf`

6.2.2 Explicit Indent Checking?

6.2.2.1 GLR Filtering

Layout-Sensitive Generalized Parsing
Erdweg S., Rendel T., Kästner C., Ostermann K. (2013) Layout-Sensitive Generalized Parsing. In: Czarnecki K., Hedin G. (eds) Software Language Engineering. SLE 2012. Lecture Notes in Computer Science, vol 7745. Springer, Berlin, Heidelberg

6.2.2.2 Data-Dependent Grammars

Ch. 3: "In this chapter we propose a parsing framework that embraces context information in its core."
`https://homepages.cwi.nl/~jurgenv/papers/PhDThesis_Ali_Afroozeh_and_Anastasia_Izmaylova.pdf`
align and offside keywords desugared to data-dependent expressions (section 3.3.5)

6.2.3   PEG Semantic Predicates

`https://gist.github.com/dmajda/04002578dd41ae8190fc`
Note that the 'I' provides an explicit token indicating that indentation is coming so this doesn't really work for arbitrary indentation like YAML.

6.2.4   Pseudo-Tokens and Pseudo-Grammars

6.2.4.1   Implicit Current Indentation

6.2.4.2   Explicit Current Indentation

6.3   Parsing Flow-Style Lists

6.4   Parsing Blank Lines in Indented Blocks

Chapter 7

Constructing a Parser

You can hand write a parser or use a parser generator to automatically generate a parser PEG is typically partially handwritten and partially generated.

7.1 Handwritten Parsers

7.2 Parser Generators

Chapter 8

A Model Data Parser

8.1   Design Choices

Appendix A

Data

A.1 Motivation

This book covers the topic of data because a parser operates on data, and there are important issues to be aware of with a common form of data called Unicode. These issues are listed at the beginning of this chapter in order to motivate the reader to explore these topics further in detail. The rest of the chapter is written to give the reader a full understanding of these issues, from the ground up, so that the reader can be confident that they have dealt with them appropriately when they arise in the context of parsing or any other context.

The first issue is that there are some symbols in the Unicode standard which have multiple definitions. This can lead to potential problems if this issue is not accounted for properly. For example, it requires a strategy for determining whether two names, for example, are the same. Another set of issues arise because it may be unclear what one means by the *length* of a sequences of symbols.

These issues with Unicode are typically either not well known or not understood by both the users and developers of parsing software. When the strategy for handling these questions is left undocumented there is risk that software will behave in an unexpected way, potentially leading to serious consequences.

A.2 Symbols

The signs, sounds, marks, and words humans use to represent ideas are called *symbols*. For example, the twenty six letters of the English alphabet are symbols representing sounds or parts of words. Words are also symbols for the ideas they represent. The numerals 0 through 9 are symbols for numbers. As long as something represents, or *symbolizes* something else, it is a symbol. We use symbols as sort of "real world" currency to remember or communicate something meaningful to ourselves or others.

A.2.1 Interpreting Symbols

Wikipedia defines *data* as any sequence of symbols for which "a specific act of interpretation" gives those symbols meaning[1]. In other words, the essential nature of data is that is in a *sequence*, which means that the symbols occur one after another, and that there is a particular meaning associated with the ordering and choice of the symbols in the data.

A.2.2 Manipulating Symbols

Computers are well known for their ability to manipulate data. A computer can read and *execute*, or follow, instructions as represented by a sequence of symbols. These instructions can, in turn, direct the manipulation of symbols stored in the computer's memory. It would be confusing to think of a computer working with the kind of visual symbols that humans typically think of like a + (plus) sign or the numeral 2, because they do not actually work directly with those kinds of symbols. The symbols that a computer can read, follow, and manipulate are quite limited. In fact at the lowest levels, there are only two symbols. Inside the computer, these symbols exist as electrical states inside the transistors and wires of the computer. In an electric model, we think of these states as *off* and *on* and by convention these states symbolize the numbers 0 and 1, respectively.

---

[1] https://en.wikipedia.org/wiki/Data_(computing)

A.3   Binary Data

A.3.1   Bits

A numerical symbol which is either 0 or 1 is known as a *binary digit*, or *bit* for short. At the most fundamental level, bits are the letters of a very simple alphabet that both humans and computers can manipulate in an equivalent way. For example, both humans and computers are capable of numerically adding together the symbols 1 and 0 to produce the symbol 1, although they go about it in different ways. It is through these common definitions and ways of processing symbols that computers can perform computations that are relevant to humans.

A.3.2   Binary Codes

Since computers are only able to manipulate sequences of bits, we have devised an incredibly rich variety of ways, known as *codes*[2] to map the information available to us into sequences of ones and zeros. For example, *decimal*[3] numbers, which are the standard numbers we use from day to day, can be converted to a *binary number* composed of binary digits, using a *binary code* as seen in Table 1.1.

Decimal	Binary
0	000
1	001
2	010
3	011
4	100
5	101
6	110
7	111

Table A.1: Map of Decimal Digits to Binary Digits

A.3.3   Bytes and Octets

A *byte* is typically 8 bits. A byte also represents the smallest unit of a computer's memory that is *addressable* by location. Computer systems store and retrieve bytes of data in specific locations of the computer's memory and these locations are identified using a number, called a *memory address*. There are examples of old computers that had more or less bits but it is standard now to assume a byte has 8 bits. In some contexts, such as with networking standards, the term *octet* may be used instead of byte if there is a concern that the definition of a byte is ambiguous.

A.3.4   Base-N Number Systems

People typically communicate using numbers which are split into digits based on the number ten. This is called the *decimal numeral system* When a digit exceeds nine, another digit to the left represents ten times the current digit. Further digits can be added to represent higher numbers based on powers of ten. Numbers in the decimal number system are also called *base-10* numbers.

With binary numbers however, when a digit exceeds *one*, another digit is added to the left that represents *twice* the current digit. Binary numbers are based on powers of two and are called *base-2* numbers.

In addition to these number systems, two other systems are commonly used with computer systems, called *hexadecimal* and *octal* numbers which are base-16 and base-8 numbers respectively.

A.3.5   Octal Notation

Octal is a numerical encoding that allows the use of 8 different symbols for each digit, 0 through 7. Each symbol encodes 3 bits of a binary number.

---

[2]https://en.wikipedia.org/wiki/Code
[3]https://en.wikipedia.org/wiki/Decimal

A.3.6 Hexadecimal Notation

Hexadecimal is a numerical encoding that allows the use of 16 different symbols for each digit, 0 through 9 and A through F. Lowercase letters are also used but mixing upper and lowercase letters is uncommon. Each symbol encodes 4 bits, or a nibble, of a binary number. Hexadecimal numbers are convenient because you can always represent one byte with only two digits.

Decimal	Binary	Octal	Hexadecimal
0	0	0	0
1	1	1	1
2	10	2	2
3	11	3	3
4	100	4	4
5	101	5	5
6	110	6	6
7	111	7	7
8	1000	10	8
9	1001	11	9
10	1010	12	A
11	1011	13	B
12	1100	14	C
13	1101	15	D
14	1110	16	E
15	1111	17	F
16	10000	18	10
17	10001	19	11
18	10010	20	12
19	10011	21	13

Table A.2: Numbers in Various Numeric Notations

A.3.7 Exercises for Section A.3

Exercise A.3.7.1: Write the numbers 8 through 15 using 4 bits.

Exercise A.3.7.2: Create a mapping from the numbers 0 through 8 to two digits of the three symbols 0, 1, and 2.

A.4 Character Encodings

Characters are symbols that represent sounds or parts of words.

A.4.1 ASCII

An early standard character encoding was ASCII, which stands for American Standard Code for Information Interchange.

A.4.2 Terminal Control Characters

A computer terminal is composed of a display monitor for displaying characters to humans and a keyboard to send characters to the computer. These characters are displayed in rows and columns.

Control characters were invented in order to control terminal operations using character codes, such as sending a character code to go to the next line. Character codes are still supported, although many of them are outdated.

A.4.3 Unicode and UTF-8

The Unicodeⓡ Standard is a universal character encoding standard for all written characters and text.

A.4.3.1   Codepoints

A.4.3.2   Graphemes

A.4.3.3   Unicode Normalization

Unicode Normalization is a process used to resolve ambiguities in the Unicode Standard.

A.5   Summary of Chapter A

Index